PROTOTYPE WALKTHROUGH

How MIRQAB Works

A step-by-step demo flow using only real screenshots of the running MIRQAB SOC Dashboard, each labelled as real, simulated or limited.

TEAM

Bdoor Fahed Alosaimi
Nouf Saleh Altalhi
Alaa Rafa Al Saad

PROJECT

MIRQAB — AI-Powered Security Operations Platform for SMBs

27 September 2026

About these screenshots

Provenance — read first

- Every image comes unmodified from the project folder `apps/soc-dashboard/e2e/screenshots/`. They were captured by the team's automated browser QA scripts (Playwright, `apps/soc-dashboard/e2e/*.cjs`) against the running application on the development machine between **19 and 26 September 2026**. The SHA-256 fingerprint of each file is listed in the appendix and matches the original.
- The images are **not** mock-ups. They show **development data**: a mix of Simulation Mode events, manual test events, public dataset events and events from the containerized Wazuh lab. **No customer or company data appears.**
- Every incident shows its source label: *Simulation (Synthetic)*, *Manual / Test*, *Open-Source Dataset* or *Live Agent*. *Live Agent* means the event arrived through a forwarder ingestion endpoint. In this development database, those events came from the lab and from test submissions.
- File names describe the QA check each image was taken for, not always the content. The captions describe what is actually visible.
- Some screenshots from 19–22 Sep show wording that was later refined when the 5-minute response window was reworked on 26 Sep 2026.
- Two demo steps (Notifications and Reports) have **no dedicated screenshot** in the project. This is stated where relevant, and nothing was substituted.

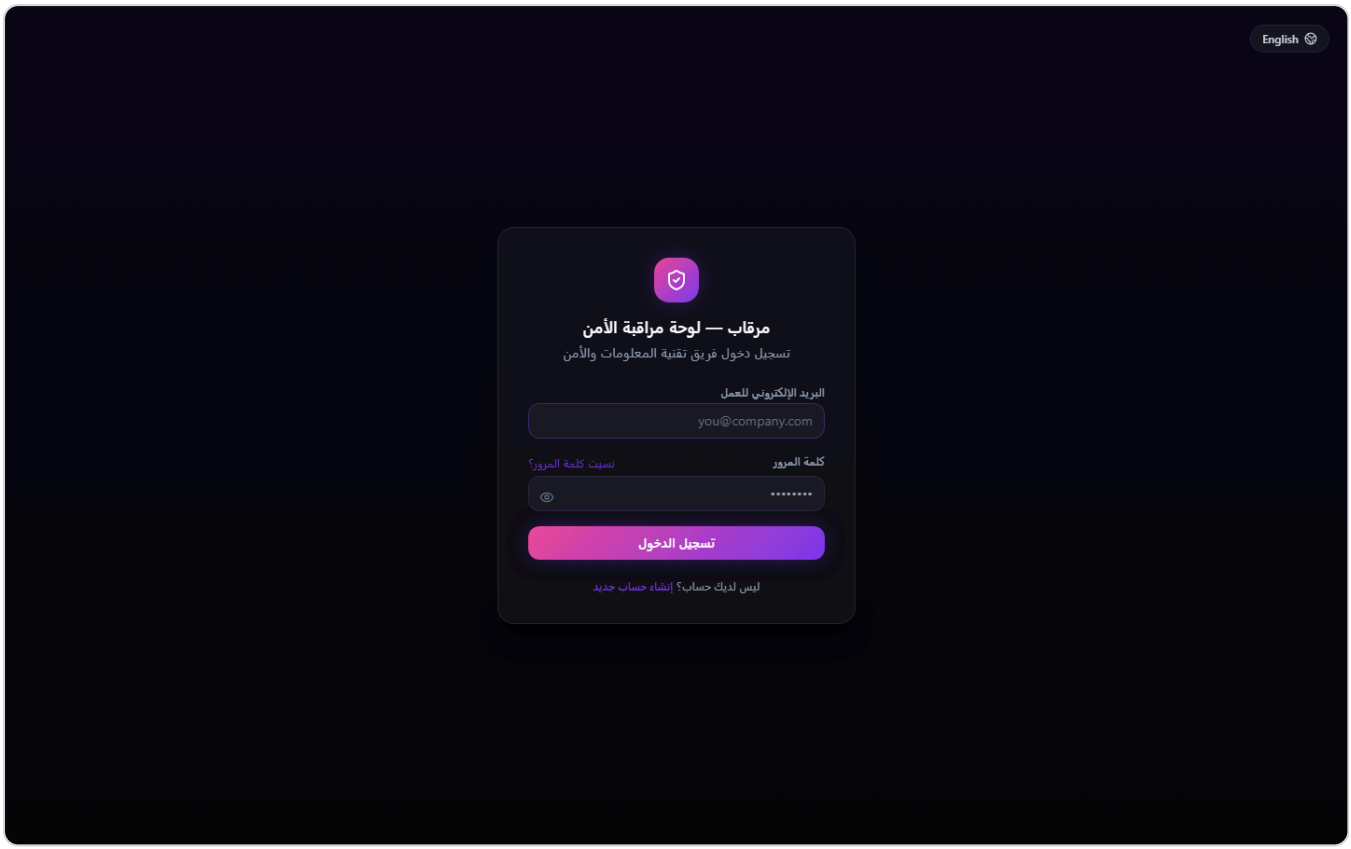
How to read the status labels

VERIFIED	Tested and confirmed: by the automated test suite (38 suites, 421 tests, all passing on 27 Sep 2026) and/or by a live test recorded in the team's engineering reports.
VERIFIED IN LAB	Confirmed end-to-end in a controlled, containerized lab. Not tested on physical employee devices.
LIMITED	Implemented and working, with a clearly stated limitation.
NOT VERIFIED	Implemented in code but not tested live.
NOT IMPLEMENTED	Not available in the application.
NOT DEPLOYED	No production deployment exists.
BLOCKED	Cannot be completed because of an external dependency.
PLANNED	A future roadmap item. Not built.
DEMO / SIMULATION	Uses synthetic or public demo data. Not a real environment.

The demo flow

#	Step	Screenshot(s)	Feature status
1	Login	01-login	VERIFIED
2	Dashboard	15-dashboard-kpi-cards-no-delta	VERIFIED
3	Incident	26-centered-two-window-en, 08-isolation-manual-pending-en	VERIFIED · data SIMULATION
4	Investigation	response-window-01-waiting-en-dark	VERIFIED
5	Evidence	06-evidence-expanded-en	VERIFIED
6	MITRE	response-timer-01-overdue-en-dark	VERIFIED LIMITED
7	Suggested Solutions	04b-inc2-suggested-solutions-en	VERIFIED
8	Response timer	response-window-01-waiting-ar-light, response-timer-02-acknowledged-en-dark	VERIFIED · automation LIMITED
9	Device / network topology	14-devices-dropdown-page (+ map in step 2)	VERIFIED
10	Isolation workflow	12-isolate-recommendation-routes-to-real-confirm, 07b-isolation-pending-approval-en	VERIFIED IN LAB
11	Notifications	*no dedicated screenshot*	VERIFIED LIMITED
12	Reports	*no dedicated screenshot*	VERIFIED

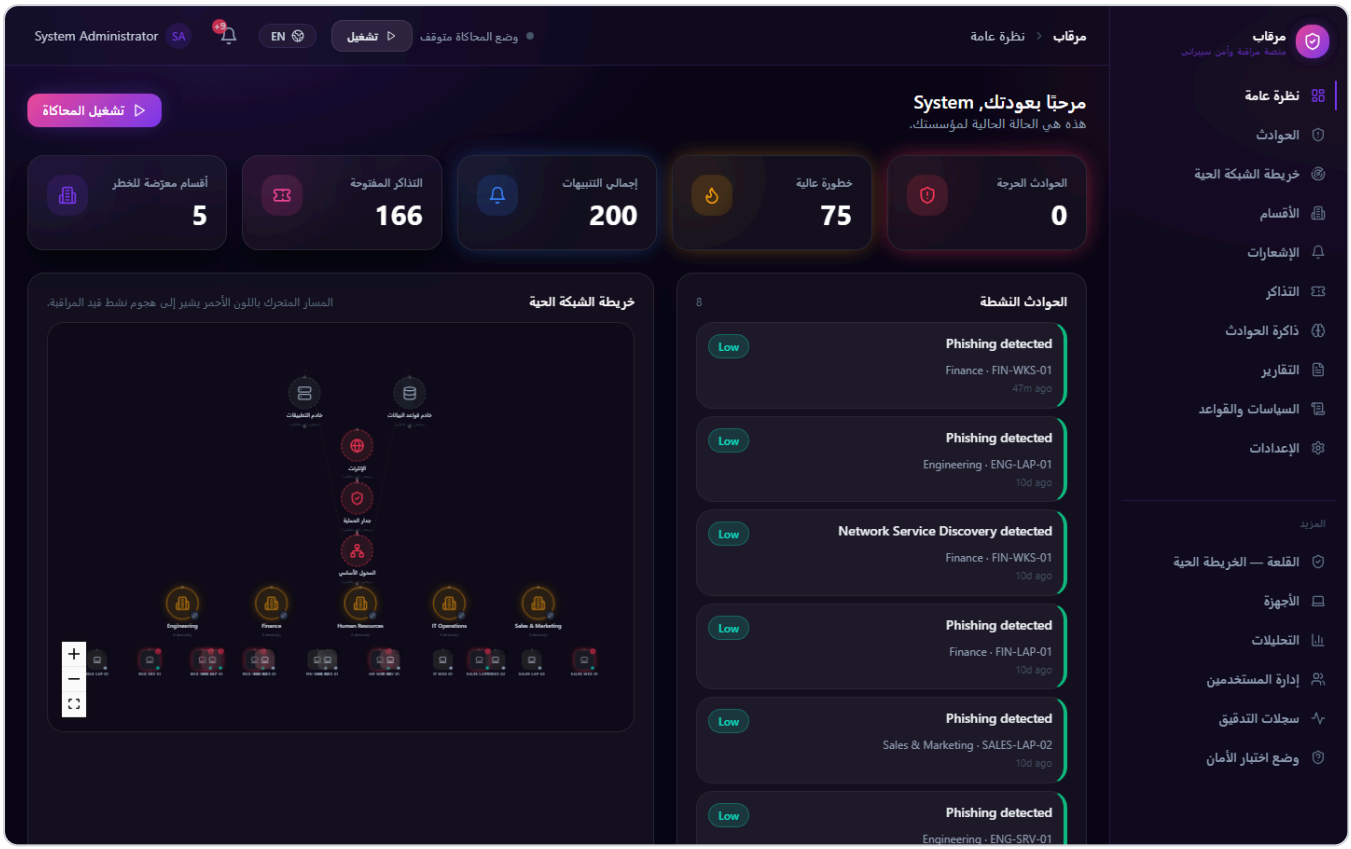
1 Login



01-login.png · captured 19 Sep 2026 · Arabic interface

What you see	The SOC Dashboard sign-in page in Arabic: "مرقاب — لوحة مراقبة الأمن" (MIRQAB — Security Monitoring Dashboard). It has a work-email field, a password field with a show/hide toggle, a "Forgot password?" link, a "Create new account" link and a language switch (top right, "English").
Feature	Authentication. Login returns a short-lived access token and a rotating refresh token. The login endpoint is rate-limited to 5 requests per minute, and wrong credentials always get the same generic message.
Status	VERIFIED — real screen of the running application.

2 Dashboard (Overview)



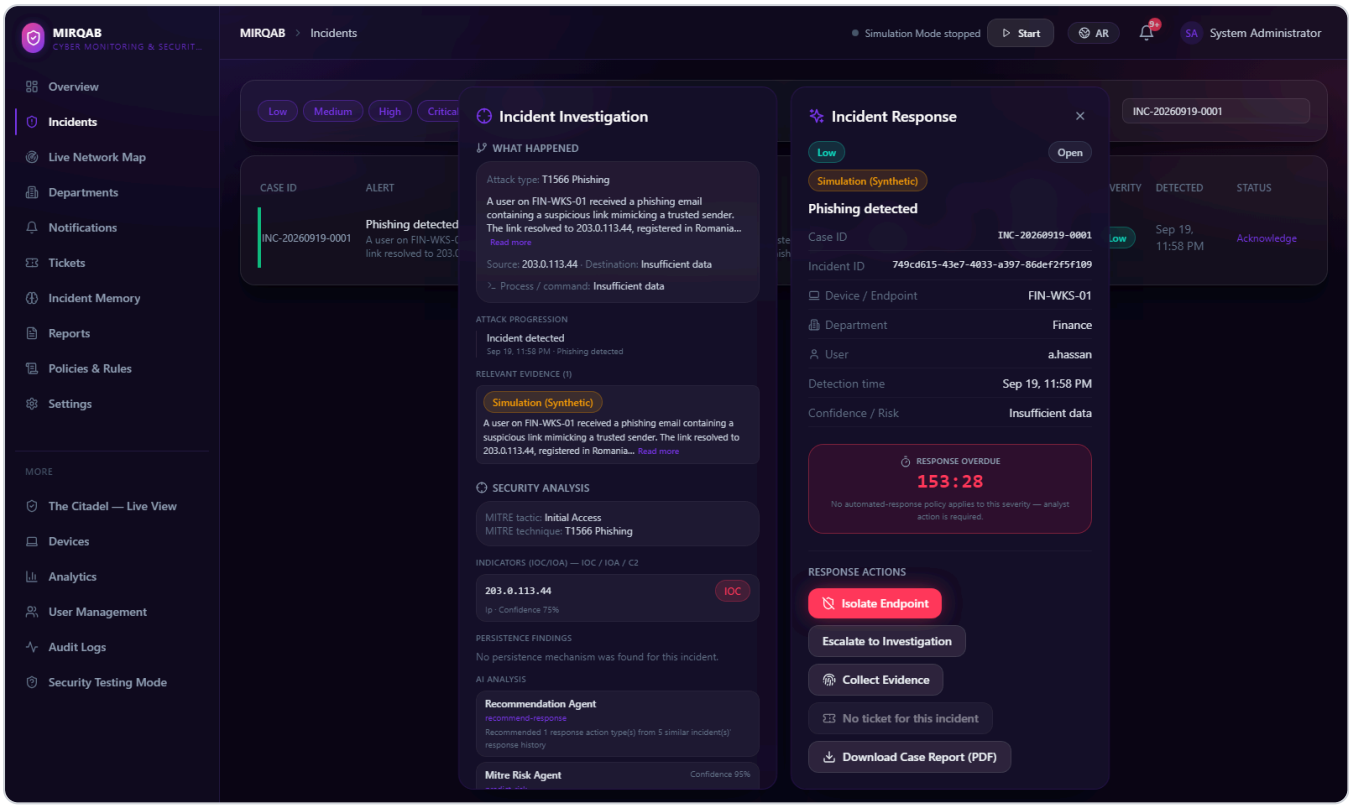
15-dashboard-kpi-cards-no-delta.png · captured 20 Sep 2026 · Arabic, dark theme

What you see The Overview page ("نظرة عامة"). Five KPI tiles: critical incidents (0), high severity (75), total alerts (200), open tickets (166), departments at risk (5). Below them are the *Live Network Map* and the *Active Incidents* list. The top bar shows the Simulation Mode status ("stopped"), a *Start* button, the language switch, the notification bell and the signed-in System Administrator.

Feature SOC overview. All figures are computed from the incident, ticket and device records in the database.

Status VERIFIED — real screen. The numbers reflect the **development database** (simulation, test, dataset and lab events accumulated during development). They are not measurements of a real organization.

3 Incident



26-centered-two-window-en.png · captured 21 Sep 2026 · English, dark theme

What you see

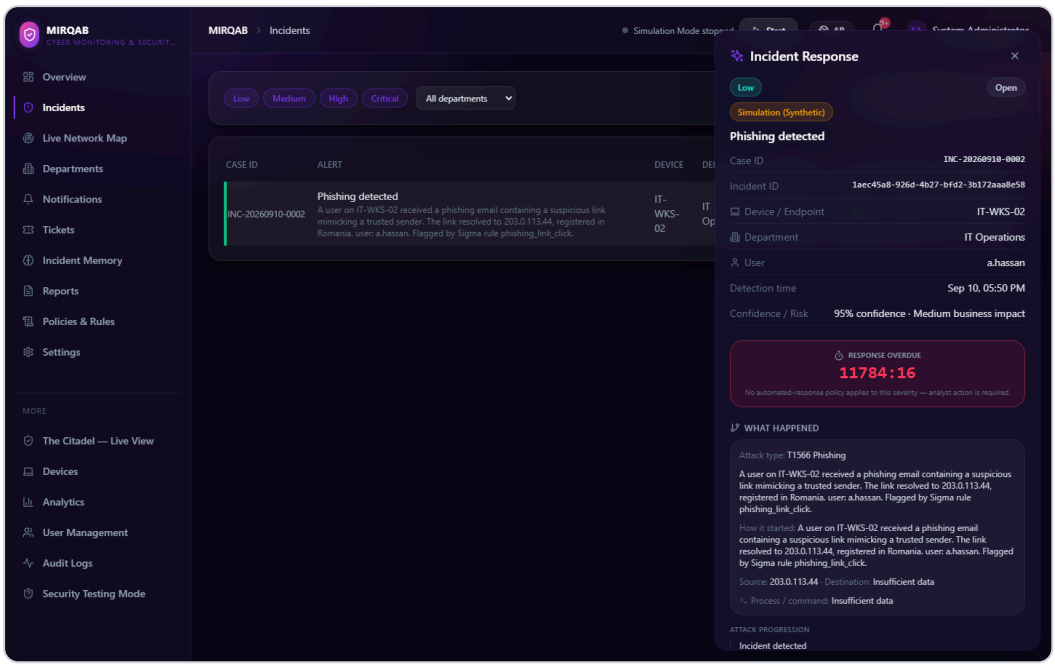
The Incidents page with case `INC-20260919-0001`, "Phishing detected", opened in the two-window incident view. **Incident Investigation** (left) shows what happened, attack progression, evidence, security analysis (MITRE tactic *Initial Access*, technique *T1566 Phishing*), the IOC `203.0.113.44` (75% confidence) and analysis entries. **Incident Response** (right) shows severity, status, Case ID, device FIN-WKS-01, department, user, a response-overdue timer and the response actions: Isolate Endpoint, Escalate to Investigation, Collect Evidence and Download Case Report (PDF).

Feature

Incident case view: correlation into a Case ID, evidence, MITRE mapping, indicators and response actions.

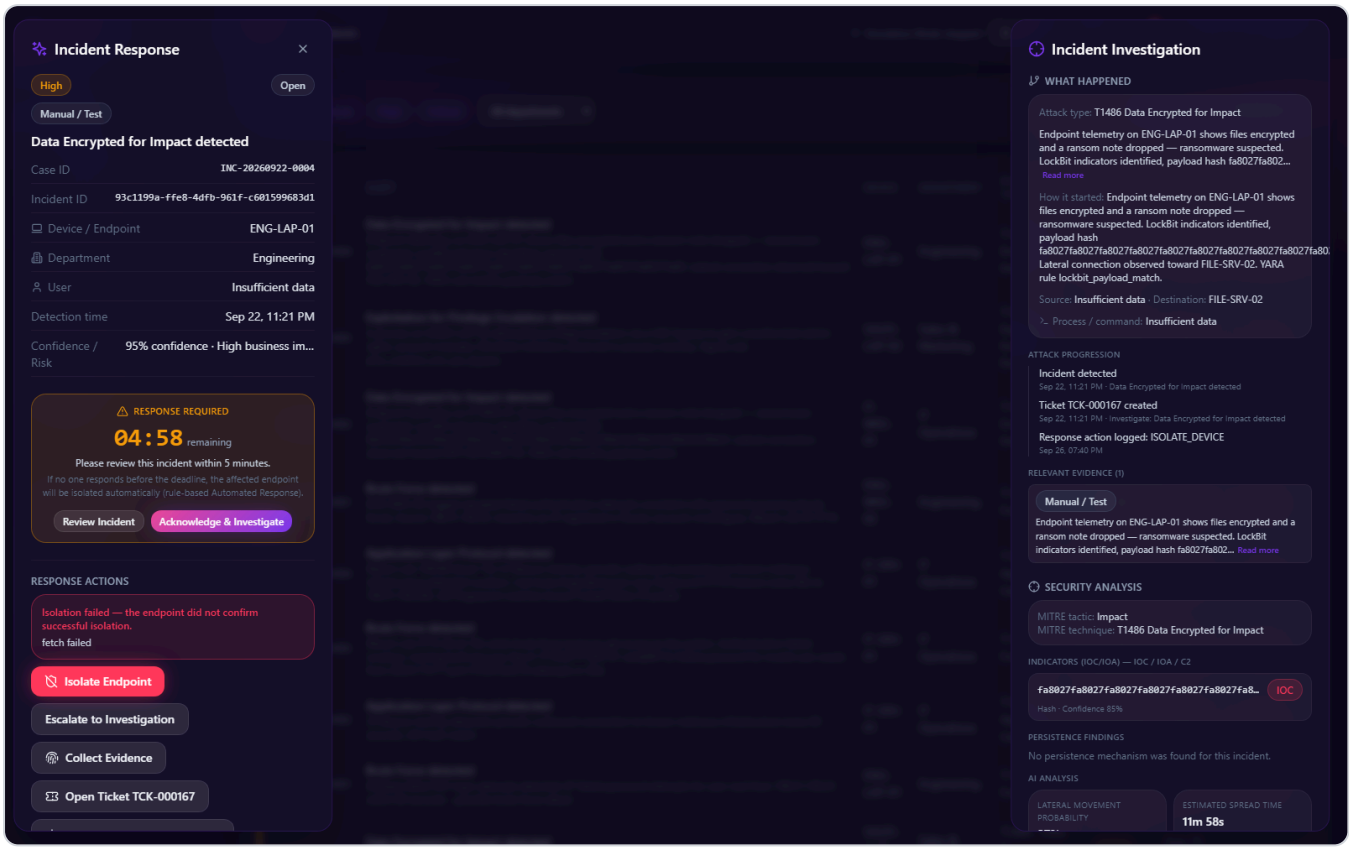
Status

VERIFIED UI. The incident itself is labelled SIMULATION (SYNTHETIC): it was generated by Simulation Mode. 203.0.113.44 is a reserved documentation IP address. This screenshot predates the 26 Sep response-window rework, so the overdue text differs slightly from the current build.



08-isolation-manual-pending-en.png · captured 19 Sep 2026 — another synthetic phishing incident, clearly labelled "Simulation (Synthetic)". For a LOW-severity incident the panel states that no automated-response policy applies and that analyst action is required.

4 Investigation



response-window-01-waiting-en-dark.png · captured 26 Sep 2026 · current response-window design

What you see

A HIGH-severity case `INC-20260922-0004`, "Data Encrypted for Impact detected" on ENG-LAP-01 (source: *Manual / Test*). The Investigation window shows the narrative (ransomware-style encryption, a payload hash, a lateral connection toward FILE-SRV-02), attack progression (incident detected → ticket TCK-000167 created → response action logged), MITRE *Impact / T1486*, a hash IOC (85% confidence) and analysis metrics (for example, lateral-movement probability). The Response window shows **RESPONSE REQUIRED — 04:58 remaining** with *Review Incident* and *Acknowledge & Investigate*, and a red notice: **"Isolation failed — the endpoint did not confirm successful isolation. fetch failed."**

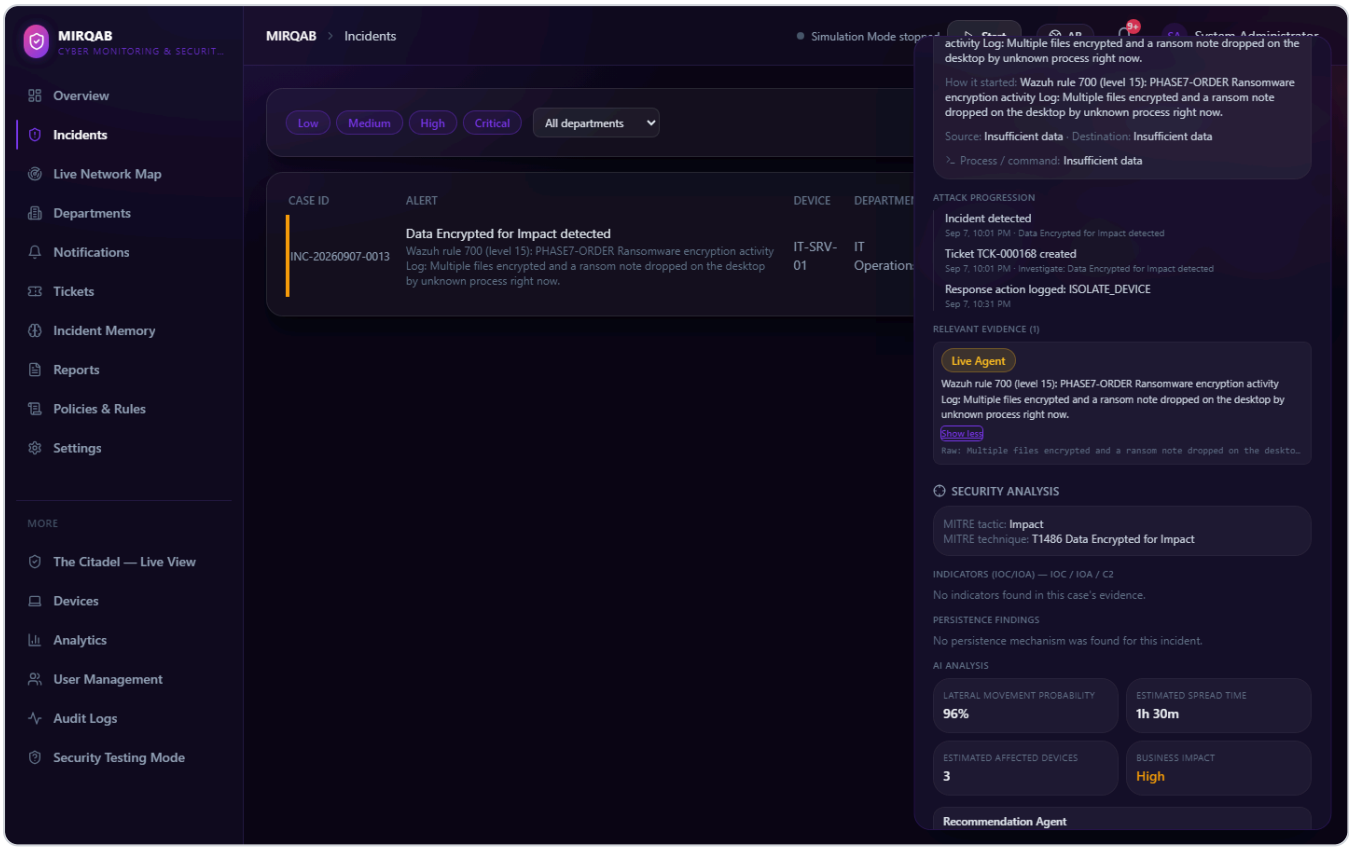
Feature

Investigation workspace and the 5-minute response window.

Status

VERIFIED UI. The incident was submitted manually for testing. The failed isolation is **genuine**: ENG-LAP-01 is a virtual device with no real endpoint behind it, so MIRQAB correctly recorded FAILED instead of pretending success. The analysis percentages are statistics from the department's past incidents, not a machine-learning prediction LIMITED.

5 Evidence



06-evidence-expanded-en.png · captured 19 Sep 2026

What you see

Case `INC-20260907-0013` on IT-SRV-01. The evidence item is labelled *Live Agent* and expanded ("Show less"), showing the full Wazuh rule text (rule 700, level 15, "PHASE7-ORDER Ransomware encryption activity") and the raw log line. Below are MITRE *Impact / T1486*, "No indicators found in this case's evidence", "No persistence mechanism was found" and four analysis tiles: lateral-movement probability 96%, estimated spread time 1 h 30 m, estimated affected devices 3, business impact High.

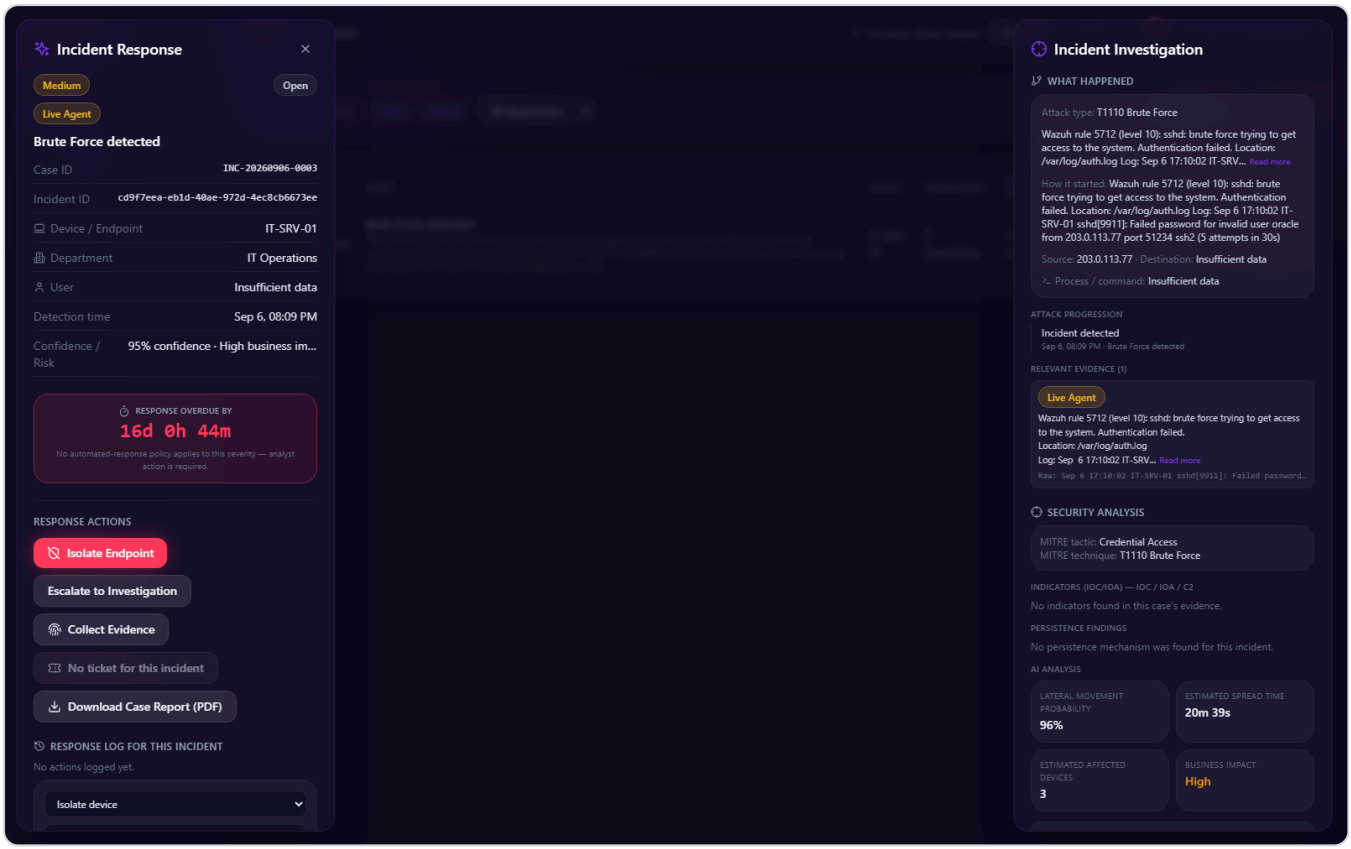
Feature

Evidence handling: each contributing event is stored with its source label, readable text and raw excerpt. Empty sections say so explicitly instead of inventing data.

Status

VERIFIED UI. The "PHASE7-ORDER" text indicates a test event from the team's security-testing phases, sent in Wazuh alert format through the ingestion endpoint. The analysis tiles are statistics, not ML predictions **LIMITED**.

6 MITRE ATT&CK mapping



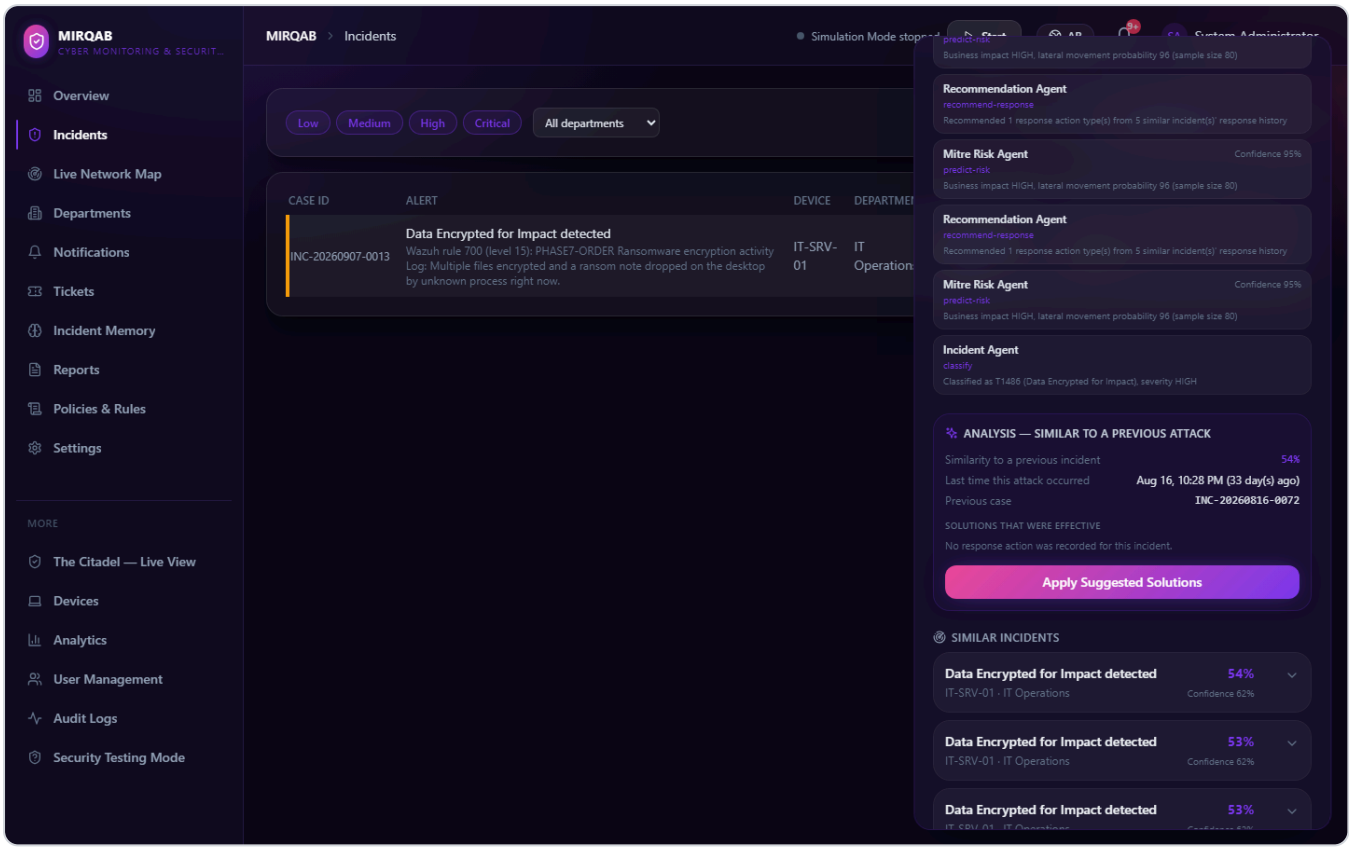
response-timer-01-overdue-en-dark.png · captured 22 Sep 2026

What you see Case `INC-20260906-0003`, "Brute Force detected" (MEDIUM) on IT-SRV-01. The evidence is the Wazuh rule 5712 sshd brute-force alert ("Failed password for invalid user oracle from 203.0.113.77 … 5 attempts in 30s"). Security Analysis maps it to MITRE tactic *Credential Access*, technique *T1110 Brute Force*. The timer shows "RESPONSE OVERDUE BY 16d 0h 44m" with "No automated-response policy applies to this severity — analyst action is required."

Feature Deterministic MITRE mapping, and the rule that MEDIUM/LOW incidents never get automated containment.

Status **VERIFIED**. Coverage is **LIMITED** to 12 curated techniques. The event came through the Wazuh ingestion endpoint during development testing (*Live Agent* label).

7 Suggested Solutions



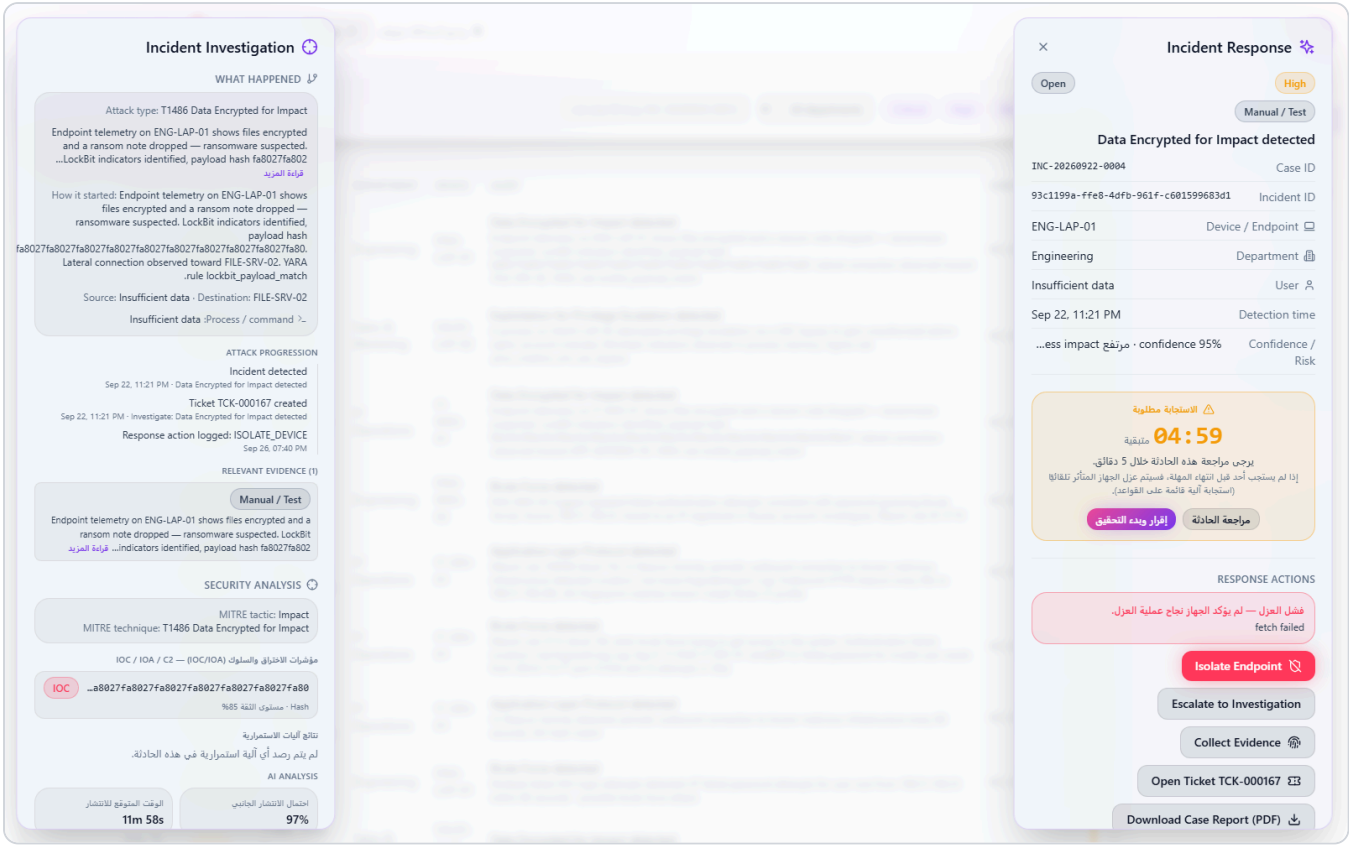
04b-inc2-suggested-solutions-en.png · captured 19 Sep 2026

What you see For case `INC-20260907-0013`, the *AI Analysis* entries (Incident Agent: "Classified as T1486 … severity HIGH"; MITRE Risk Agent: business impact and lateral-movement statistics with sample size; Recommendation Agent: "Recommended 1 response action type(s) from 5 similar incident(s)' response history"). Below is **Analysis — Similar to a previous attack**: 54% similarity to `INC-20260816-0072` from 33 days earlier, "Solutions that were effective", the **Apply Suggested Solutions** button, and a ranked **Similar Incidents** list (54%, 53%, 53%).

Feature Incident memory, similar-incident search and Suggested Solutions. Suggestions come from a technique-specific fix catalog plus actions recorded on similar past incidents. High-impact actions require analyst approval.

Status VERIFIED. The "agents" are named, least-privilege wrappers around **deterministic** services. No language model was running (see the AI architecture section of the Project Submission).

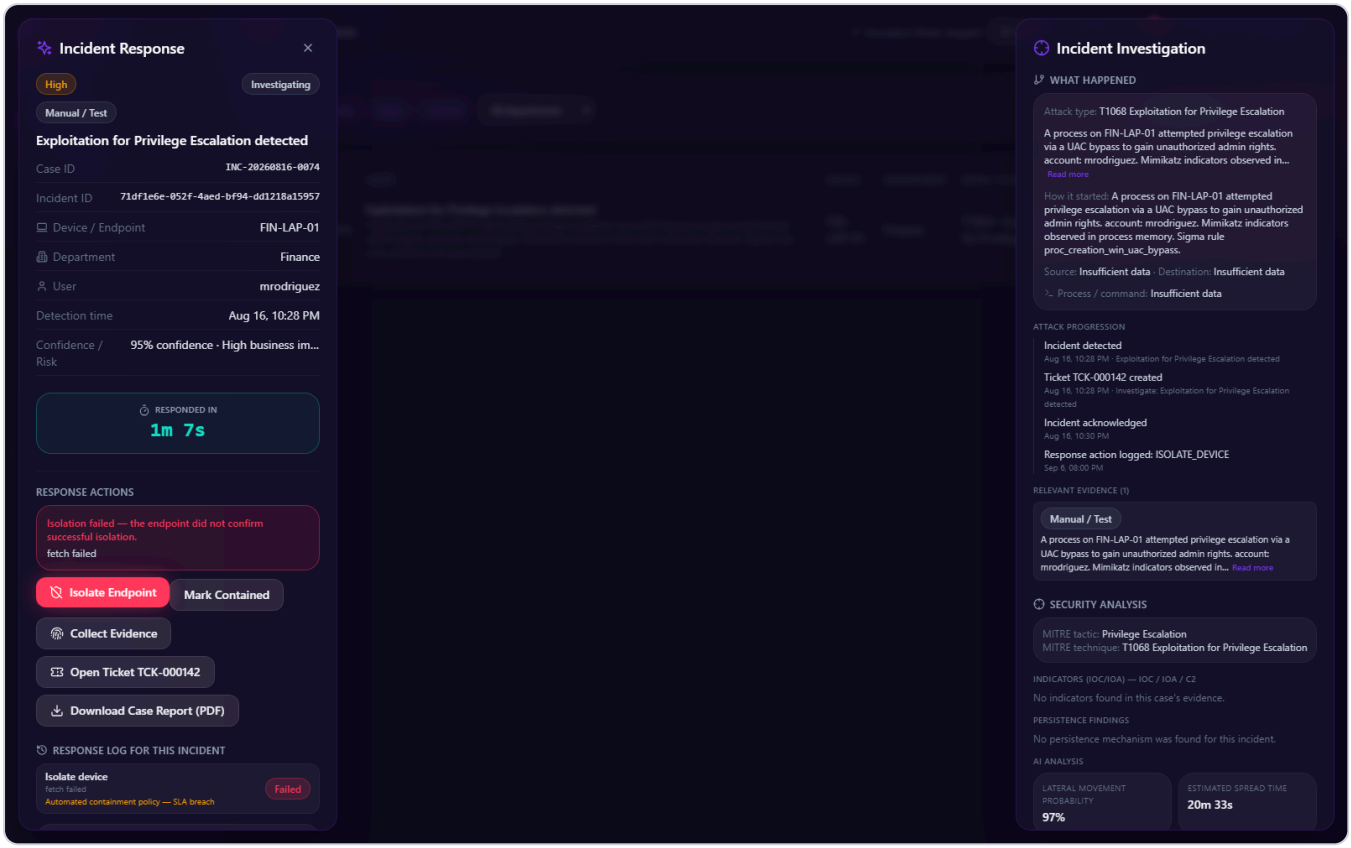
8 Response timer



response-window-01-waiting-ar-light.png · captured 26 Sep 2026 · Arabic, light theme

What you see The same HIGH incident as step 4, in Arabic with the light theme. The response box reads "الاستجابة مطلوبة — 04:59 متبقية" (Response required — 04:59 remaining): review within 5 minutes, otherwise the endpoint will be isolated automatically by the rule-based automated response. It has "Review Incident" and "Acknowledge & Investigate" buttons.

Status VERIFIED UI. LIMITED: several panel labels remain in English in Arabic mode.



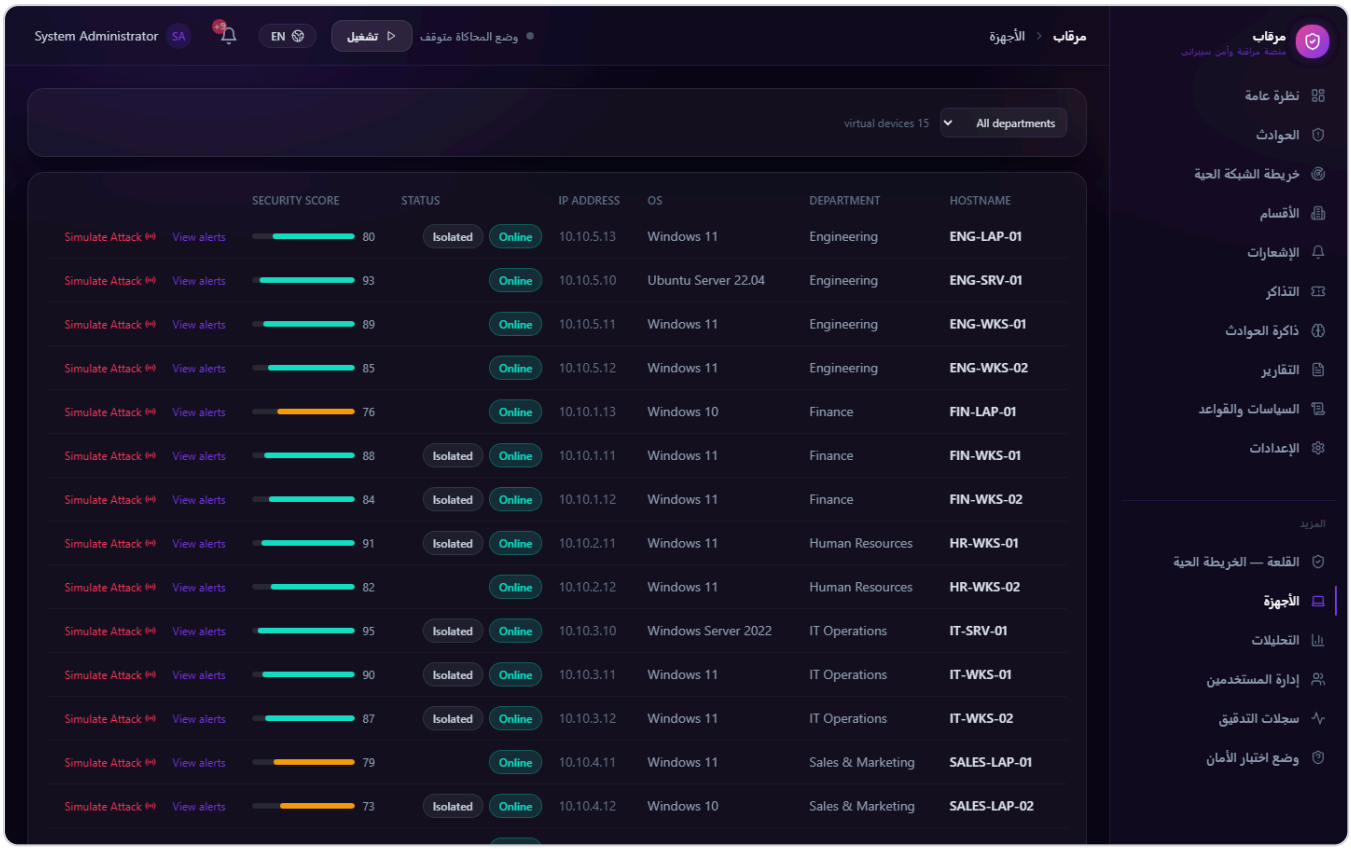
response-timer-02-acknowledged-en-dark.png · captured 22 Sep 2026

What you see A HIGH incident (`INC-20260816-0074` , T1068 privilege escalation, status *Investigating*) whose timer shows **RESPONDED IN 1m 7s**, which means it was acknowledged inside the window. The response log shows an *Isolate device* entry marked **Failed** ("fetch failed — Automated containment policy — SLA breach"), plus *Mark Contained*, *Collect Evidence*, the ticket link and *Download Case Report (PDF)*.

Feature Acknowledgement, and the honest recording of response outcomes.

Status VERIFIED workflow. The failed automated isolation shows the LIMITED state of automated containment in development: the attempt was made and correctly recorded as FAILED because no reachable endpoint existed. The log entry appears to come from the earlier automated-containment policy that the 26 Sep rework replaced.

9 Device and network topology



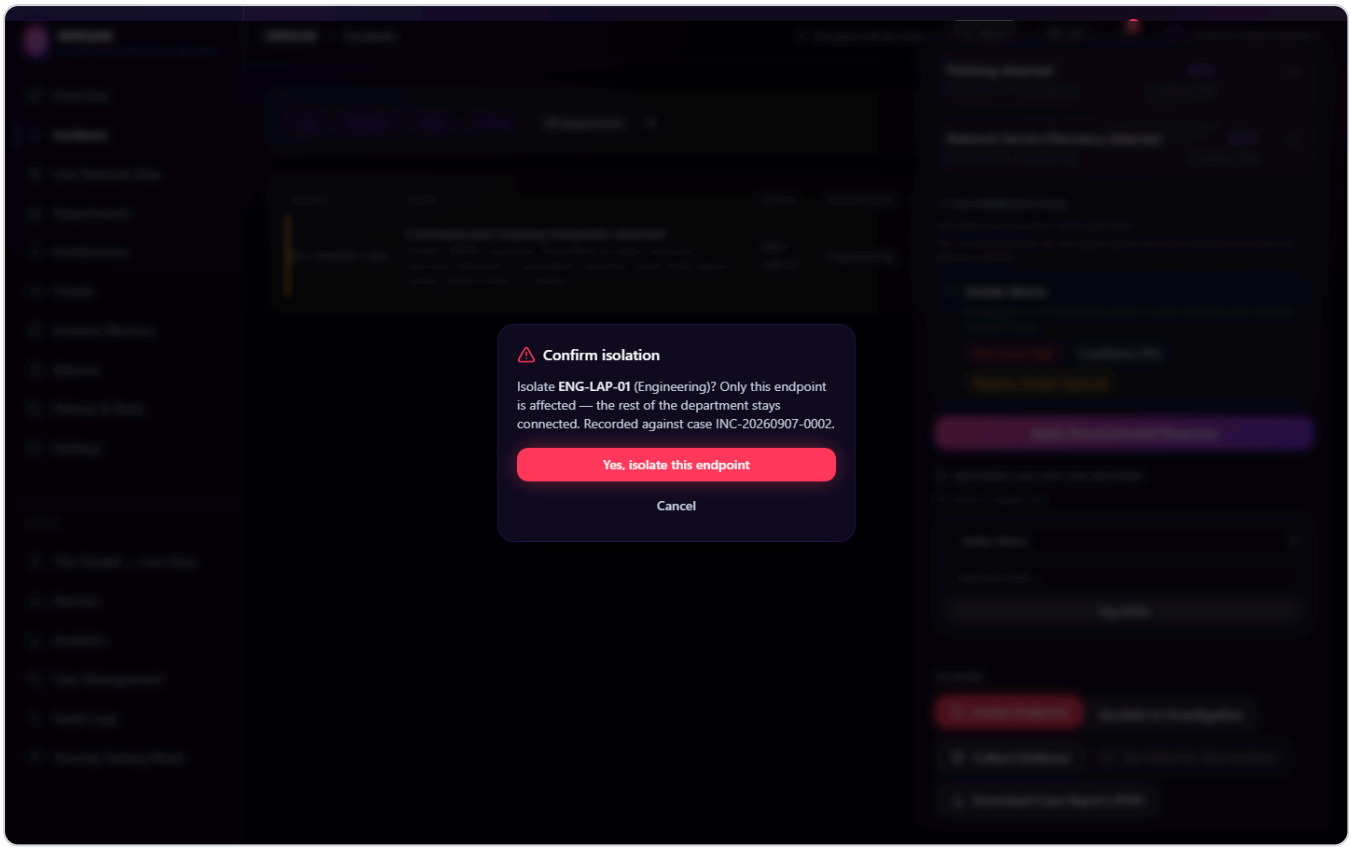
14-devices-dropdown-page.png · captured 20 Sep 2026 · Arabic interface

What you see — The Devices page ("الأجهزة"), labelled "virtual devices 15", with a department filter. Each row shows hostname, department, OS, IP address, status (Online, and Isolated where applicable), a security score bar, "View alerts" and "Simulate Attack".

Feature — Device inventory, which feeds the live network map (visible in step 2: internet → firewall → core switch → departments → devices, with active attacks highlighted).

Status — VERIFIED UI. DEMO: these are the 15 seeded **virtual** devices. The "Isolated" badges reflect isolation actions recorded as successful in the development database during testing. **No physical employee device was isolated.**

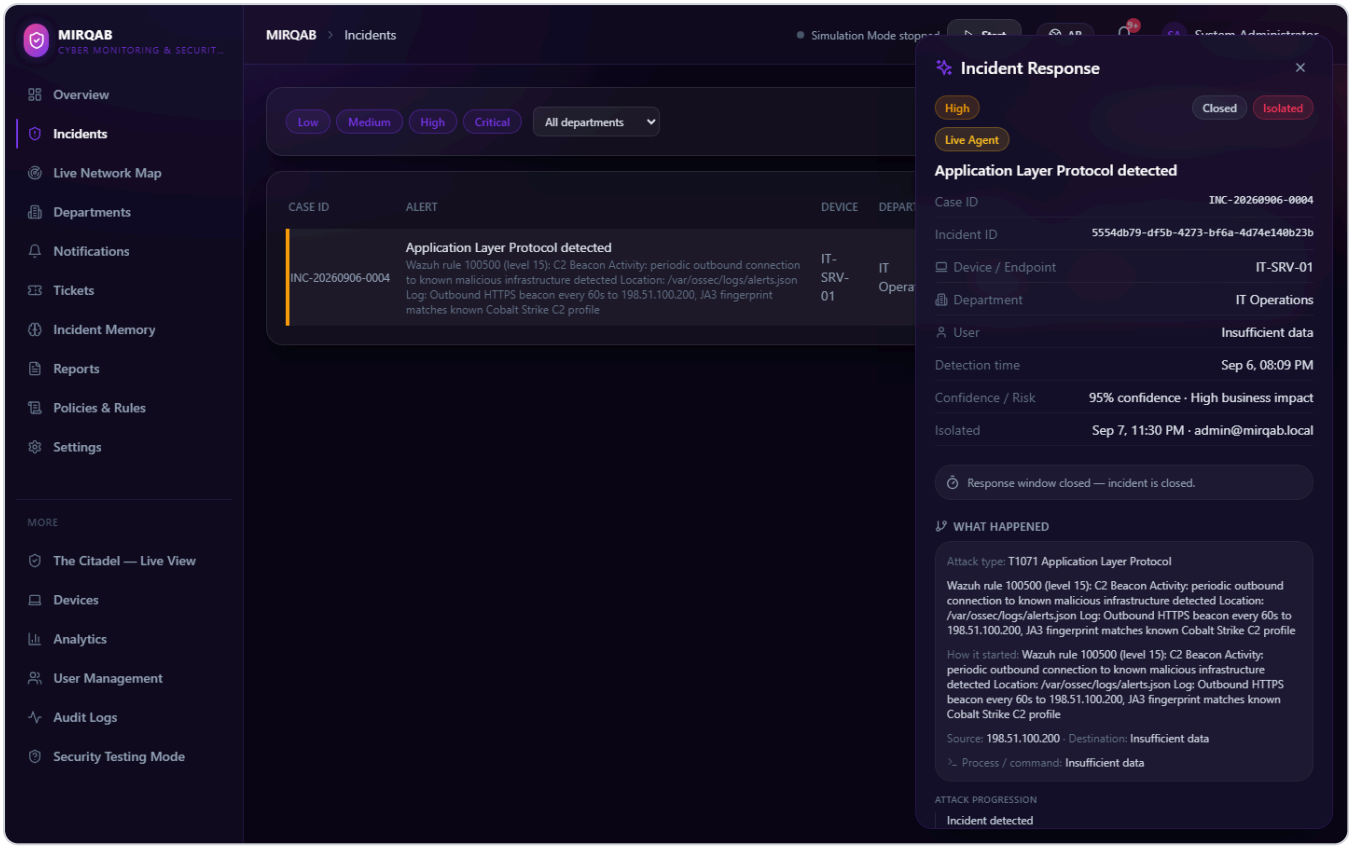
10 Isolation workflow



12-isolate-recommendation-routes-to-real-confirm.png · captured 19 Sep 2026

What you see The **Confirm isolation** dialog: "Isolate ENG-LAP-01 (Engineering)? Only this endpoint is affected — the rest of the department stays connected. Recorded against case INC-20260907-0002." with "Yes, isolate this endpoint" and "Cancel".

Feature A human confirmation step before isolation. Accepting an isolation recommendation leads to this same confirmation; nothing is executed silently.



07b-isolation-pending-approval-en.png · captured 19 Sep 2026 (the file name refers to the QA check, not the content)

What you see Case `INC-20260906-0004`, "Application Layer Protocol detected" (HIGH, C2 beacon evidence, MITRE T1071) on **IT-SRV-01**, status **Closed** with an **Isolated** badge: "Isolated — Sep 7, 11:30 PM · admin@mirqab.local". The response window reads "closed — incident is closed".

Feature A completed isolation. The device shows "Isolated" only when a genuine SUCCESS isolation action exists.

Status VERIFIED IN LAB. IT-SRV-01 is the name of the containerized **safe-test Wazuh agent** used in the lab isolation tests (security Phase 2A: 2/2 successful end-to-end isolations, real firewall change confirmed by agent callback). No physical device was involved. Releasing isolation from MIRQAB is NOT IMPLEMENTED.

11 Notifications

No dedicated screenshot available

The project's screenshot folder contains no capture of the Notifications page, and none was created for this document. The notification **bell with its unread counter** ("9+") is visible in the top bar of the screenshots above (for example steps 3 and 7).

How it works: new HIGH/CRITICAL incidents create a "response required" notification, and an expired response window creates a "containment result" notification. They are delivered in real time over Socket.IO to exactly the SOC users allowed to see the incident: Admin, Security Analyst, and the IT Technician of that department. The Notifications page lists them with filters (incident, ticket, policy, audit).

Status: VERIFIED by unit tests. LIMITED: there is no per-employee notification, because the backend has no employee-ownership model for incidents.

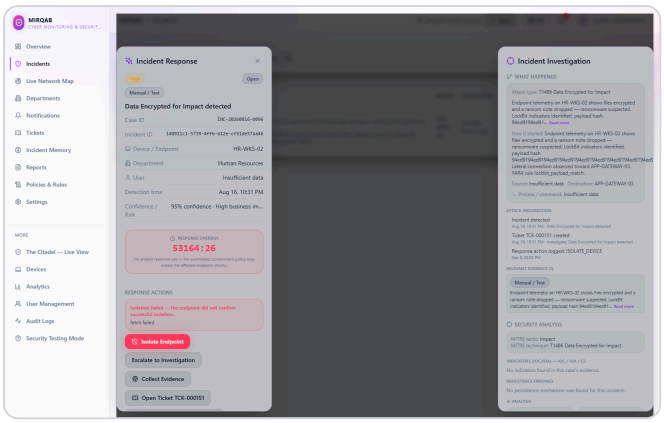
12. Reports

No dedicated screenshot available

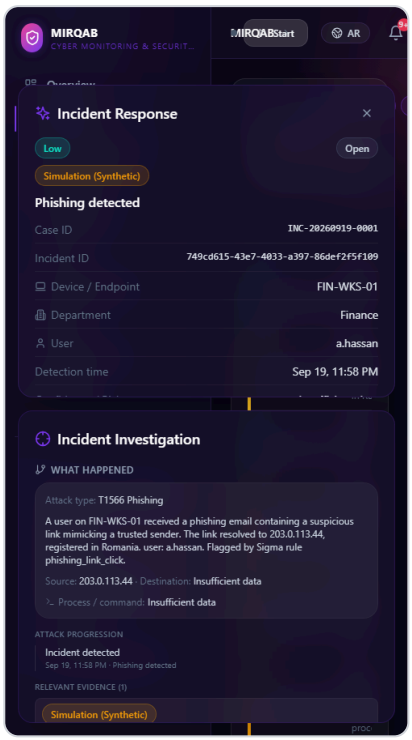
The screenshot folder contains no capture of a generated report or of the Reports page. The **Download Case Report (PDF)** action is visible in the Incident Response window (steps 3, 4 and 8).

How it works: the case report PDF has 11 sections (summary, evidence, timeline, IOC, IOA, persistence, C2, AI analysis, response actions, recommendations, audit trail). The weekly SOC report is generated every Monday 05:00 UTC for the just-completed week (or on demand) and is downloadable from the Reports page. **Status:** VERIFIED by unit tests.

+ Themes, languages and mobile



backdrop-fix-03-en-light.png · 22 Sep 2026 — light theme, English: HIGH ransomware test incident on HR-WKS-02 (*Manual / Test*), with a failed isolation recorded as "fetch failed". The overdue wording predates the 26 Sep rework.



21-two-window-stacked-mobile.png · 21 Sep 2026 — narrow (mobile-width) layout: the two incident windows stack vertically. **LIMITED**: the header items overlap slightly at this width.

Appendix — screenshot fingerprints (SHA-256)

Each file in `02_Prototype/screenshots/` is byte-identical to the original in `apps/soc-dashboard/e2e/screenshots/`.

File	SHA-256
01-login.png	`5ae66f85a229c1ffba89d2d33b20018443a17b4eded496e8883f17f3197920be`
04b-inc2-suggested-solutions-en.png	`54dde04f5a6aacfec77f381b0c6b8d3cc5f142b2747e1cf685d1b4f339617852`
06-evidence-expanded-en.png	`1226e2b2e41fb260b14f2476e7c07bb7e4e14453b3e7355474108b010177c612`
07b-isolation-pending-approval-en.png	`91718ac6e772f8e29c38389ac55088b005565697aa98afa679544134ae21d713`
08-isolation-manual-pending-en.png	`b15de592fd4083eeed03b0c34f942846bc9a885ddf7f5f897c1913398756caaa`
12-isolate-recommendation-routes-to-real-confirm.png	`38bcd916b34734f014a96038b9088bb1f7819686c8a8aa901bce33999f576893`
14-devices-dropdown-page.png	`2dd74b140eea5e2e85b7aef2c80b5fcbd1ac4cc2b2ce00d0ec6facb6c1064a19`
15-dashboard-kpi-cards-no-delta.png	`1af404dd804ee3446ded0a44082b71b677aee96aa53fcd67f5019814e3c8bf6e`
21-two-window-stacked-mobile.png	`ce914ced4c5335733389a776462337759d46fa05bc72c2c9e6871d88bb815bca`
26-centered-two-window-en.png	`5f3da230d257be93e50d9354c7137facac934eaab84217588f17c117134532f5`
backdrop-fix-03-en-light.png	`840008fde5dd96319557e042966be6130420e09db7ba1e455edd9d2c2154beb2`

File	SHA-256
response-timer-01-overdue-en-dark.png	`f76821a428b30b037506680c7887bc9293d0ff25c58be95eaf262643e519c4e6`
response-timer-02-acknowledged-en-dark.png	`bca0e13098b38af6edfbc1fa3b4f0f2e36a48f268c4fc4bc1e39996711003b8a`
response-window-01-waiting-ar-light.png	`806ce10eb32f53629b2bd6fb2cf376ca230af487d6fbadb94ec9f37bc505a693`
response-window-01-waiting-en-dark.png	`e54f99d89d4df9dd8f807fad8a1d064dabf1efedd23015fbdecb897aa8ba71d8`